

```
/// *****  
>>> start-all.sh - 하둡 데몬들 실행시킵니다.  
>>> jps - 실행중인 하둡 데몬들(자바 프로세서) 확인합니다.  
/// *****
```

```
chosun@chosun-VirtualBox:~/hadoop$ start-all.sh  
starting namenode, logging to /home/chosun/hadoop-1.2.1/libexec/../logs/hadoop-chosun-namenode-chosun-VirtualBox.out  
WARNING: An illegal reflective access operation has occurred  
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()  
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil  
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
WARNING: All illegal access operations will be denied in a future release  
localhost: starting datanode, logging to /home/chosun/hadoop-1.2.1/libexec/../logs/hadoop-chosun-datanode-chosun-VirtualBox.out  
localhost: WARNING: An illegal reflective access operation has occurred  
localhost: WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()  
localhost: WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil  
localhost: WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
localhost: WARNING: All illegal access operations will be denied in a future release  
localhost: starting secondarynamenode, logging to /home/chosun/hadoop-1.2.1/libexec/../logs/hadoop-chosun-secondarynamenode-chosun-VirtualBox.out  
localhost: WARNING: An illegal reflective access operation has occurred  
localhost: WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()  
localhost: WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil  
localhost: WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
localhost: WARNING: All illegal access operations will be denied in a future release  
starting jobtracker, logging to /home/chosun/hadoop-1.2.1/libexec/../logs/hadoop-chosun-jobtracker-chosun-VirtualBox.out  
WARNING: An illegal reflective access operation has occurred  
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()  
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil  
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
WARNING: All illegal access operations will be denied in a future release  
localhost: starting tasktracker, logging to /home/chosun/hadoop-1.2.1/libexec/../logs/hadoop-chosun-tasktracker-chosun-VirtualBox.out  
localhost: WARNING: An illegal reflective access operation has occurred  
localhost: WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()  
localhost: WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil  
localhost: WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
localhost: WARNING: All illegal access operations will be denied in a future release  
chosun@chosun-VirtualBox:~/hadoop$
```

```
chosun@chosun-VirtualBox:~/hadoop$ jps  
15510 Jps  
14855 DataNode  
15080 SecondaryNameNode  
15209 JobTracker  
15420 TaskTracker  
4286 org.eclipse.equinox.launcher_1.5.300.v20190213-1655.jar  
14670 NameNode  
chosun@chosun-VirtualBox:~/hadoop$
```

```
/// *****  
>>> HWEx7_2MapSideJoin.jar은 맵 사이드 조인을 수행하는 jar파일입니다.  
>>> HWEx7_3ReduceSideJoin.jar은 리듀스 사이드 조인을 수행하는 jar파일입니다.  
/// *****
```

```
chosun@chosun-VirtualBox:~/hadoop$ ls  
1987.csv  HWEx7_3ReduceSideJoin.jar  HomeWorkEX6_2_Seq.jar  README.txt  conf  hadoop-ant-1.2.1.jar  
hadoop-minicluster-1.2.1.jar  ivy.xml  pids  test.py  2008.csv  HomeWorkEX6_1.jar  HomeWorkEX6_3_Total.jar  bin  contrib  
hadoop-client-1.2.1.jar  hadoop-test-1.2.1.jar  lib  sbin  webapps  CHANGES.txt  HomeWorkEX6_2_Map.jar  LICENSE.txt  
build.xml  dataexpo  hadoop-core-1.2.1.jar  hadoop-tools-1.2.1.jar  libexec  share  HWEx7_2MapSideJoin.jar  
HomeWorkEX6_2_Search.jar  NOTICE.txt  c++  docs  hadoop-examples-1.2.1.jar  ivy  logs  src
```



```
chosun@chosun-VirtualBox:~/hadoop$ hadoop fs -ls
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
Found 5 items
drwxr-xr-x - chosun supergroup 0 2019-05-12 13:28 /user/chosun/HW_0512
drwxr-xr-x - chosun supergroup 0 2019-05-13 19:58 /user/chosun/HW_0513_meta
drwxr-xr-x - chosun supergroup 0 2019-05-12 14:09 /user/chosun/HW_Mapfile
drwxr-xr-x - chosun supergroup 0 2019-05-12 14:13 /user/chosun/HW_Sequencefile
drwxr-xr-x - chosun supergroup 0 2019-05-12 14:14 /user/chosun/HW_total
chosun@chosun-VirtualBox:~/hadoop$
chosun@chosun-VirtualBox:~/hadoop$ hadoop fs -put carriers_new2.csv HW_0513_meta/carriers_new.csv
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
chosun@chosun-VirtualBox:~/hadoop$
chosun@chosun-VirtualBox:~/hadoop$ hadoop fs -ls HW_0513_meta
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
Found 1 items
-rw-r--r-- 1 chosun supergroup 43740 2019-05-13 19:59 /user/chosun/HW_0513_meta/carriers_new.csv
chosun@chosun-VirtualBox:~/hadoop$
/// ***** ///
>>> [Ex7.2]
>>> jar 파일을 실행합니다.
>>> HW_0513_meta/carriers_new.csv와 HW_0512를 처리한 것을 map_join디렉토리에 저장합니다.
>>> MapSideJoin은 클래스 이름입니다.
/// ***** ///
chosun@chosun-VirtualBox:~/hadoop$ hadoop jar HWEx7_2MapSideJoin.jar MapSideJoin HW_0513_meta/carriers_new.csv HW_0512 map_join
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
19/05/13 20:06:04 INFO input.FileInputFormat: Total input paths to process : 1
19/05/13 20:06:04 INFO util.NativeCodeLoader: Loaded the native-hadoop library
19/05/13 20:06:04 WARN snappy.LoadSnappy: Snappy native library not loaded
19/05/13 20:06:05 INFO mapred.JobClient: Running job: job_201905131956_0001 // job 시작
19/05/13 20:06:06 INFO mapred.JobClient: map 0% reduce 0% // 맵 시작, 리듀스 --- x
19/05/13 20:06:16 INFO mapred.JobClient: map 18% reduce 0% // ~~ ing
19/05/13 20:06:21 INFO mapred.JobClient: map 36% reduce 0%
19/05/13 20:06:27 INFO mapred.JobClient: map 54% reduce 0%
19/05/13 20:06:32 INFO mapred.JobClient: map 72% reduce 0%
19/05/13 20:06:37 INFO mapred.JobClient: map 90% reduce 0% // ~~ ing
19/05/13 20:06:40 INFO mapred.JobClient: map 100% reduce 0% // 맵-리듀스 끝
19/05/13 20:06:41 INFO mapred.JobClient: Job complete: job_201905131956_0001 // job 끝
19/05/13 20:06:41 INFO mapred.JobClient: Counters: 19
19/05/13 20:06:41 INFO mapred.JobClient: Map-Reduce Framework
19/05/13 20:06:41 INFO mapred.JobClient: Spilled Records=0 // 작업 중에 디스크에 기록 된 총 레코드 수
19/05/13 20:06:41 INFO mapred.JobClient: Virtual memory (bytes) snapshot=22930333696 // 모든 작업에 사용된 총 가상 메모리
```

```
19/05/13 20:06:41 INFO mapred.JobClient: Map input records=7009728 // 맵에 입력되는 데이터의 수
19/05/13 20:06:41 INFO mapred.JobClient: SPLIT_RAW_BYTES=1265 // 맵에서 가져온 입력 분할 개체의 바이트 수
19/05/13 20:06:41 INFO mapred.JobClient: Map output records=7009728 // 맵의 출력으로 생성된 데이터의 수
19/05/13 20:06:41 INFO mapred.JobClient: Physical memory (bytes) snapshot=1690693632 // 모든 작업에 사용되는 총 물리적 메모리
19/05/13 20:06:41 INFO mapred.JobClient: CPU time spent (ms)=35830 // 소요된 cpu시간
19/05/13 20:06:41 INFO mapred.JobClient: Total committed heap usage (bytes)=961544192 // JVM에 사용할 수 있는 총 메모리 양
19/05/13 20:06:41 INFO mapred.JobClient: File Input Format Counters
19/05/13 20:06:41 INFO mapred.JobClient: Bytes Read=689454004 // 모든 파일 시스템에 대해 모든 태스크에서 읽은 데이터 양
19/05/13 20:06:41 INFO mapred.JobClient: FileSystemCounters
19/05/13 20:06:41 INFO mapred.JobClient: HDFS_BYTES_READ=689455269 // HDFS에서 읽은 데이터 양
19/05/13 20:06:41 INFO mapred.JobClient: FILE_BYTES_WRITTEN=642775 // 로컬 파일 시스템에 기록된 데이터 양
19/05/13 20:06:41 INFO mapred.JobClient: HDFS_BYTES_WRITTEN=900468672 // HDFS에 기록된 데이터 양
19/05/13 20:06:41 INFO mapred.JobClient: Job Counters
19/05/13 20:06:41 INFO mapred.JobClient: Launched map tasks=11 // 실행된 총 맵 작업 수
19/05/13 20:06:41 INFO mapred.JobClient: SLOTS_MILLIS_REDUCES=0 // 모든 리듀스 작업에 의해 소요된 총 시간(ms)
19/05/13 20:06:41 INFO mapred.JobClient: Total time spent by all reduces waiting after reserving slots (ms)=0
19/05/13 20:06:41 INFO mapred.JobClient: SLOTS_MILLIS_MAPS=51496 // 모든 맵 작업에 의해 사용한 총 시간(ms)
19/05/13 20:06:41 INFO mapred.JobClient: Total time spent by all maps waiting after reserving slots (ms)=0
19/05/13 20:06:41 INFO mapred.JobClient: Data-local map tasks=11 // 필요한 데이터가 포함된 노드에서 시작된 맵 태스크 수
19/05/13 20:06:41 INFO mapred.JobClient: File Output Format Counters
19/05/13 20:06:41 INFO mapred.JobClient: Bytes Written=900468672 // 모든 파일 시스템에 대해 모든 태스크에서 작성된 데이터 양
```

MR-Job Result:0

chosun@chosun-VirtualBox:~/hadoop\$

/// ***** ///

>>> 실행 결과 :: 리듀서가 없는 맵리듀스 잡이므로 매퍼의 출력 데이터가 모두 생성되었습니다.
>>> 만약 리듀서를 정의했다면 part-m-xxxx파일이 아니라, 리듀서의 출력 데이터인 part-r-xxxx파일들이 생성됩니다.
>>> 첫 번째 출력 데이터인 part-m-00000을 조회해본 결과,
항공사 이름과 해당 항공사에 대한 운항 통계 데이터가 정상적으로 저장돼 있음을 확인했습니다.

/// ***** ///

chosun@chosun-VirtualBox:~/hadoop\$ **hadoop fs -ls map_join**

WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release

```
Found 13 items
-rw-r--r--  1 chosun supergroup      0 2019-05-13 20:06 /user/chosun/map_join/_SUCCESS
drwxr-xr-x  - chosun supergroup      0 2019-05-13 20:06 /user/chosun/map_join/_logs
-rw-r--r--  1 chosun supergroup 87401259 2019-05-13 20:06 /user/chosun/map_join/part-m-00000
-rw-r--r--  1 chosun supergroup 87045872 2019-05-13 20:06 /user/chosun/map_join/part-m-00001
-rw-r--r--  1 chosun supergroup 88157694 2019-05-13 20:06 /user/chosun/map_join/part-m-00002
-rw-r--r--  1 chosun supergroup 88954716 2019-05-13 20:06 /user/chosun/map_join/part-m-00003
-rw-r--r--  1 chosun supergroup 87780517 2019-05-13 20:06 /user/chosun/map_join/part-m-00004
-rw-r--r--  1 chosun supergroup 87481529 2019-05-13 20:06 /user/chosun/map_join/part-m-00005
-rw-r--r--  1 chosun supergroup 87195853 2019-05-13 20:06 /user/chosun/map_join/part-m-00006
-rw-r--r--  1 chosun supergroup 87122067 2019-05-13 20:06 /user/chosun/map_join/part-m-00007
-rw-r--r--  1 chosun supergroup 86756011 2019-05-13 20:06 /user/chosun/map_join/part-m-00008
-rw-r--r--  1 chosun supergroup 89365213 2019-05-13 20:06 /user/chosun/map_join/part-m-00009
-rw-r--r--  1 chosun supergroup 23207941 2019-05-13 20:06 /user/chosun/map_join/part-m-00010
```

chosun@chosun-VirtualBox:~/hadoop\$

chosun@chosun-VirtualBox:~/hadoop\$ **hadoop fs -cat map_join/part-m-00000 | head -10**

WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release

```
WN      Southwest Airlines Co.2008,1,3,4,2003,1955,2211,2225,WN,335,N712SW,128,150,116,-14,8,IAD,TPA,810,4,8,0,,0,NA,NA,NA,NA,NA
WN      Southwest Airlines Co.2008,1,3,4,754,735,1002,1000,WN,3231,N772SW,128,145,113,2,19,IAD,TPA,810,5,10,0,,0,NA,NA,NA,NA,NA
WN      Southwest Airlines Co.2008,1,3,4,628,620,804,750,WN,448,N428WN,96,90,76,14,8,IND,BWI,515,3,17,0,,0,NA,NA,NA,NA,NA
```

```
WN      Southwest Airlines Co.2008,1,3,4,926,930,1054,1100,WN,1746,N612SW,88,90,78,-6,-4,IND,BWI,515,3,7,0,,0,NA,NA,NA,NA,NA
WN      Southwest Airlines Co.2008,1,3,4,1829,1755,1959,1925,WN,3920,N464WN,90,90,77,34,34,IND,BWI,515,3,10,0,,0,2,0,0,0,32
WN      Southwest Airlines Co.2008,1,3,4,1940,1915,2121,2110,WN,378,N726SW,101,115,87,11,25,IND,JAX,688,4,10,0,,0,NA,NA,NA,NA,NA
WN      Southwest Airlines Co.2008,1,3,4,1937,1830,2037,1940,WN,509,N763SW,240,250,230,57,67,IND,LAS,1591,3,7,0,,0,10,0,0,0,47
WN      Southwest Airlines Co.2008,1,3,4,1039,1040,1132,1150,WN,535,N428WN,233,250,219,-18,-1,IND,LAS,1591,7,7,0,,0,NA,NA,NA,NA,NA
WN      Southwest Airlines Co.2008,1,3,4,617,615,652,650,WN,11,N689SW,95,95,70,2,2,IND,MCI,451,6,19,0,,0,NA,NA,NA,NA,NA
WN      Southwest Airlines Co.2008,1,3,4,1620,1620,1639,1655,WN,810,N648SW,79,95,70,-16,0,IND,MCI,451,3,6,0,,0,NA,NA,NA,NA,NA
cat: Unable to write to output stream.
chosun@chosun-VirtualBox:~/hadoop$
/// ***** ///
>>> [Ex7.3]
>>> jar 파일을 실행합니다.
>>> HW_0513_meta/carriers_new.csv와 HW_0512를 처리한 것을 reduce_join_output 디렉토리에 저장합니다.
>>> ReduceSideJoin은 클래스 이름입니다.
/// ***** ///
chosun@chosun-VirtualBox:~/hadoop$ hadoop jar HWEx7_3ReduceSideJoin.jar ReduceSideJoin HW_0513_meta/carriers_new.csv HW_0512
reduce_join_output
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/
hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
19/05/13 20:10:50 INFO input.FileInputFormat: Total input paths to process : 1
19/05/13 20:10:50 INFO util.NativeCodeLoader: Loaded the native-hadoop library
19/05/13 20:10:50 WARN snappy.LoadSnappy: Snappy native library not loaded
19/05/13 20:10:50 INFO input.FileInputFormat: Total input paths to process : 1
19/05/13 20:10:50 INFO mapred.JobClient: Running job: job_201905131956_0002 // job 시작
19/05/13 20:10:51 INFO mapred.JobClient: map 0% reduce 0% // 맵-리듀스 시작
19/05/13 20:11:09 INFO mapred.JobClient: map 1% reduce 0% // ~~ ing
19/05/13 20:11:11 INFO mapred.JobClient: map 9% reduce 0%
19/05/13 20:11:12 INFO mapred.JobClient: map 16% reduce 0%
19/05/13 20:11:19 INFO mapred.JobClient: map 33% reduce 0%
19/05/13 20:11:20 INFO mapred.JobClient: map 33% reduce 2%
19/05/13 20:11:23 INFO mapred.JobClient: map 33% reduce 11%
19/05/13 20:11:27 INFO mapred.JobClient: map 41% reduce 11%
19/05/13 20:11:28 INFO mapred.JobClient: map 50% reduce 11%
19/05/13 20:11:35 INFO mapred.JobClient: map 66% reduce 16%
19/05/13 20:11:44 INFO mapred.JobClient: map 83% reduce 19%
19/05/13 20:11:47 INFO mapred.JobClient: map 91% reduce 22%
19/05/13 20:11:48 INFO mapred.JobClient: map 100% reduce 22%
19/05/13 20:11:50 INFO mapred.JobClient: map 100% reduce 33%
19/05/13 20:11:53 INFO mapred.JobClient: map 100% reduce 70%
19/05/13 20:11:56 INFO mapred.JobClient: map 100% reduce 73%
19/05/13 20:11:59 INFO mapred.JobClient: map 100% reduce 77%
19/05/13 20:12:02 INFO mapred.JobClient: map 100% reduce 80%
19/05/13 20:12:05 INFO mapred.JobClient: map 100% reduce 83%
19/05/13 20:12:08 INFO mapred.JobClient: map 100% reduce 86%
19/05/13 20:12:11 INFO mapred.JobClient: map 100% reduce 89%
19/05/13 20:12:14 INFO mapred.JobClient: map 100% reduce 92%
19/05/13 20:12:17 INFO mapred.JobClient: map 100% reduce 96%
19/05/13 20:12:20 INFO mapred.JobClient: map 100% reduce 99% // ~~ ing
19/05/13 20:12:22 INFO mapred.JobClient: map 100% reduce 100% // 맵-리듀스 끝
19/05/13 20:12:22 INFO mapred.JobClient: Job complete: job_201905131956_0002 // job 끝
19/05/13 20:12:22 INFO mapred.JobClient: Counters: 29
19/05/13 20:12:22 INFO mapred.JobClient: Map-Reduce Framework
19/05/13 20:12:22 INFO mapred.JobClient: Spilled Records=20662279 // 작업 중에 디스크에 기록 된 총 레코드 수
19/05/13 20:12:22 INFO mapred.JobClient: Map output materialized bytes=773579277 // 실제로 디스크에 기록된 맵 출력 바이트 수
19/05/13 20:12:22 INFO mapred.JobClient: Reduce input records=7011219 // 리듀스에 입력된 데이터의 수
19/05/13 20:12:22 INFO mapred.JobClient: Virtual memory (bytes) snapshot=27140116480 // 모든 작업에 사용된 총 가상 메모리
19/05/13 20:12:22 INFO mapred.JobClient: Map input records=7011219 // 맵에 입력되는 데이터의 수
```

```
19/05/13 20:12:22 INFO mapred.JobClient: SPLIT_RAW_BYTES=2995 // 맵에서 가져온 입력 분할 개체의 바이트 수
19/05/13 20:12:22 INFO mapred.JobClient: Map output bytes=759556767 // 모든 맵에 의해 생성된 압축되지 않은 출력 바이트수
19/05/13 20:12:22 INFO mapred.JobClient: Reduce shuffle bytes=773579277 // 리듀스 단계에 셔플된 맵 출력의 바이트 수
19/05/13 20:12:22 INFO mapred.JobClient: Physical memory (bytes) snapshot=3207630848 // 모든 작업에 사용되는 총 물리적 메모리
19/05/13 20:12:22 INFO mapred.JobClient: Reduce input groups=1491 // 리듀스 단계에서 처리한 개별 키 그룹의 수
19/05/13 20:12:22 INFO mapred.JobClient: Combine output records=0 // 출력 레코드 수
19/05/13 20:12:22 INFO mapred.JobClient: Reduce output records=7009708 // 리듀스의 출력으로 생성된 데이터의 수
19/05/13 20:12:22 INFO mapred.JobClient: Map output records=7011219 // 맵의 출력으로 생성된 데이터의 수
19/05/13 20:12:22 INFO mapred.JobClient: Combine input records=0 // 입력 레코드 수
19/05/13 20:12:22 INFO mapred.JobClient: CPU time spent (ms)=87640 // 소요된 cpu시간
19/05/13 20:12:22 INFO mapred.JobClient: Total committed heap usage (bytes)=2125463552 // JVM에 사용할 수 있는 총 메모리 양
19/05/13 20:12:22 INFO mapred.JobClient: File Input Format Counters
19/05/13 20:12:22 INFO mapred.JobClient: Bytes Read=0 // 모든 파일 시스템에 대해 모든 태스크에서 읽은 데이터 양
19/05/13 20:12:22 INFO mapred.JobClient: FileSystemCounters
19/05/13 20:12:22 INFO mapred.JobClient: HDFS_BYTES_READ=689494775 // HDFS에서 읽은 데이터 양
19/05/13 20:12:22 INFO mapred.JobClient: FILE_BYTES_WRITTEN=2280337398 // 로컬 파일 시스템에 기록된 데이터 양
19/05/13 20:12:22 INFO mapred.JobClient: FILE_BYTES_READ=1505990302 // 로컬 파일 시스템에서 읽은 데이터 양
19/05/13 20:12:22 INFO mapred.JobClient: HDFS_BYTES_WRITTEN=900468672 // HDFS에 기록된 데이터 양
19/05/13 20:12:22 INFO mapred.JobClient: Job Counters
19/05/13 20:12:22 INFO mapred.JobClient: Launched map tasks=12 // 실행된 총 맵 작업 수
19/05/13 20:12:22 INFO mapred.JobClient: Launched reduce tasks=1 // 실행된 총 리듀스 작업 수
19/05/13 20:12:22 INFO mapred.JobClient: SLOTS_MILLIS_REDUCES=48340 // 모든 리듀스 작업에 의해 소요된 총 시간(ms)
19/05/13 20:12:22 INFO mapred.JobClient: Total time spent by all reduces waiting after reserving slots (ms)=0
19/05/13 20:12:22 INFO mapred.JobClient: SLOTS_MILLIS_MAPS=57017 // 모든 맵 작업에 의해 사용한 총 시간(ms)
19/05/13 20:12:22 INFO mapred.JobClient: Total time spent by all maps waiting after reserving slots (ms)=0
19/05/13 20:12:22 INFO mapred.JobClient: Data-local map tasks=12 // 필요한 데이터가 포함된 노드에서 시작된 맵 태스크 수
19/05/13 20:12:22 INFO mapred.JobClient: File Output Format Counters
19/05/13 20:12:22 INFO mapred.JobClient: Bytes Written=900468672 // 모든 파일 시스템에 대해 모든 태스크에서 작성된 데이터 양
MR-Job Result:0
```

chosun@chosun-VirtualBox:~/hadoop\$

/// ***** ///

>>> 출력 경로를 조회해본 결과, 로그 파일과 출력 데이터가 생성되었음을 확인했습니다.

/// ***** ///

chosun@chosun-VirtualBox:~/hadoop\$ **hadoop fs -ls reduce_join_output**

WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release

```
Found 3 items
-rw-r--r--  1 chosun supergroup      0 2019-05-13 20:12 /user/chosun/reduce_join_output/_SUCCESS
drwxr-xr-x  - chosun supergroup      0 2019-05-13 20:10 /user/chosun/reduce_join_output/_logs
-rw-r--r--  1 chosun supergroup 900468672 2019-05-13 20:11 /user/chosun/reduce_join_output/part-r-00000
```

chosun@chosun-VirtualBox:~/hadoop\$

/// ***** ///

>>> 저장된 텍스트 파일을 조회해본 결과, 항공사 이름과 항공 운항 통계 데이터가 정상적으로 출력된 것을 확인했습니다.

/// ***** ///

chosun@chosun-VirtualBox:~/hadoop\$ **hadoop fs -cat reduce_join_output/part-r-00000 | head -10**

WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/home/chosun/hadoop-1.2.1/hadoop-core-1.2.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release

```
9E      Pinnacle Airlines Inc.    2008,12,26,5,1418,1250,1549,1432,9E,5860,91819E,91,102,66,77,88,ATL,PIT,526,3,22,0,,0,0,0,77,0,0
9E      Pinnacle Airlines Inc.    2008,12,28,7,1408,1250,1541,1432,9E,5860,96049E,93,102,65,69,78,ATL,PIT,526,8,20,0,,0,69,0,0,0,0
9E      Pinnacle Airlines Inc.2008,12,29,1,1247,1250,1418,1432,9E,5860,96019E,91,102,72,-14,-3,ATL,PIT,526,6,13,0,,0,NA,NA,NA,NA,NA
9E      Pinnacle Airlines Inc.2008,12,30,2,1245,1250,1419,1432,9E,5860,91669E,94,102,69,-13,-5,ATL,PIT,526,5,20,0,,0,NA,NA,NA,NA,NA
9E      Pinnacle Airlines Inc.    2008,12,3,3,1453,1440,1547,1530,9E,5860,91819E,54,50,30,17,13,ATL,TYS,152,2,22,0,,0,0,0,4,0,13
9E      Pinnacle Airlines Inc.    2008,12,5,5,1448,1440,1604,1530,9E,5860,91479E,76,50,29,34,8,ATL,TYS,152,5,42,0,,0,0,0,26,0,8
```

9E Pinnacle Airlines Inc. 2008,12,1,1,1156,1049,1511,1355,9E,5860,96049E,135,126,109,76,67,AUS,ATL,813,9,17,0,,0,0,0,9,0,67
9E Pinnacle Airlines Inc. 2008,12,2,2,1115,1049,1431,1355,9E,5860,91879E,136,126,117,36,26,AUS,ATL,813,9,10,0,,0,0,0,10,0,26
9E Pinnacle Airlines Inc. 2008,12,3,3,1054,1049,1413,1355,9E,5860,91819E,139,126,110,18,5,AUS,ATL,813,15,14,0,,0,0,0,18,0,0
9E Pinnacle Airlines Inc. 2008,12,4,4,1057,1049,1426,1355,9E,5860,96059E,149,126,120,31,8,AUS,ATL,813,12,17,0,,0,8,0,23,0,0

cat: Unable to write to output stream.

chosun@chosun-VirtualBox:~/hadoop\$

/// ***** ///

>>> 실행중인 하둡 데몬 종료합니다.

/// ***** ///

chosun@chosun-VirtualBox:~/hadoop\$ stop-all.sh

stopping jobtracker

localhost: stopping tasktracker

stopping namenode

localhost: stopping datanode

localhost: stopping secondarynamenode

chosun@chosun-VirtualBox:~/hadoop\$

chosun@chosun-VirtualBox:~/hadoop\$ jps

14467 Jps

4286 org.eclipse.equinox.launcher_1.5.300.v20190213-1655.jar

chosun@chosun-VirtualBox:~/hadoop\$

chosun@chosun-VirtualBox:~/hadoop\$

/// ***** ///

[정리]

조인을 단독으로 처리하는 맵 사이트 방식에 비해 리듀스 사이트 방식이 번거로워 보이지만,
조인할 데이터의 규모가 클 경우 리듀스 사이트 방식으로 조인을 수행해야 합니다.

/// ***** ///